



UNIVERSITY OF AMSTERDAM

BACHELOR THESIS

Conspiracy videos on YouTube

April 24, 2021

Student:
Roan Schellingerhout

Supervisor:
Dr. Maarten Marx

1 Introduction

1 YouTube attracts an average of 34.6 billion page views per month, making
2 it the world's largest video-sharing website and the second largest website on
3 the entire internet (Neufeld, 2021). The overwhelming majority of those page
4 views come from users watching videos, 70% of which are recommended to
5 users by YouTube's algorithm (Cooper, 2020). All types of content get pro-
6 duced and consumed on the website. However, conspiracy content has been
7 booming on YouTube (Donzelli et al., 2018). Alt-right (or far-right) and con-
8 spiracy channels are starting to grow their audiences, which could have many
9 negative consequences for society at large. For example, the number of people
10 who are distrustful of science is increasing, a development in which conspiracy
11 content on YouTube plays a role. Whenever this increased distrust relates to
12 important topics, such as believing in the efficacy of vaccines, it can create gen-
13 uine dangers to the public. As it turns out, more than half of the American
14 population has doubts about - or is definitely against - taking the COVID-19
15 vaccine (Rosenbaum, 2021). To better understand how YouTube's algorithm
16 (and recommender algorithms in general) allow(s) conspiracy content to thrive,
17 this research will investigate how quickly the algorithm develops a preference
18 for conspiracy video; in other words: how many videos a user needs to watch
19 before they get sent *down the rabbit hole*.

20 This research takes inspiration from the Dutch television program
21 *Zondag met Lubach*, where a similar idea was executed on a smaller scale
22 (Lubach, 2020). Lubach's experiment consisted of creating a new YouTube ac-
23 count on a never-used laptop, after which the account was used to watch a few
24 (recommended) conspiracy videos, to see how this would affect the YouTube
25 homepage of the account. After a mere three videos, the homepage of the
26 account was filled to the brim with conspiracy content, mostly having to do
27 with the coronavirus. The extremely interesting results that came from this

28 micro-experiment formed the motivation to research this phenomenon more ex-
29 tensively.

1.1 Research question

30 To research this subject, the following research question has been formulated:
31 *What is the impact of different watch strategies on the number of conspiracy*
32 *videos that have to be watched until a user's YouTube-recommendations start*
33 *preferring conspiracy content?* In this scenario, 'preferring' will be defined as
34 the situation in which the amount of conspiracy videos present in the recom-
35 mendations is significantly higher than that of the baseline.

36 In order to answer the research question, three sub-questions will have
37 to be answered. These questions are the following:

- 38 • How do different watch strategies on YouTube influence the type of content
39 that is recommended to a user?
- 40 • How long does it take for a YouTube recommendations to stop preferring
41 conspiracy videos, once they have started doing so?
- 42 • What type of classifier performs the best when it comes to labeling con-
43 spiracy videos on YouTube?

2 Theoretical Framework

2.1 Filter bubbles on social media

44 Whenever the user of a website finds themselves in their own information uni-
45 verse, in which the content and recommendations play into the user's preexisting
46 opinions and believes, they are in a filter bubble (Pariser, 2011). Users are by
47 themselves in such bubbles and each bubble is unique. Different bubbles can
48 have overlap, but each bubble is precisely tuned to an individual. In traditional
49 media, a user makes a conscious *choice* what types of opinions they want to
50 hear, for example by choosing to watch a broadcaster with a specific political
51 opinion. Online this decision is implicit: based on the user's behavior, their
52 content is filtered automatically by an algorithm, without explicit consent.

2.2 Filter bubbles on YouTube

53 Previous research has found that YouTube's recommendation algorithm runs
54 the risk of creating filter bubbles. Roth et al. (2020) came to this conclusion
55 after they analysed YouTube recommendations based on content. YouTube has
56 two distinct types of recommendations: recommendations based on the user's
57 viewing behavior and recommendations based on the content of the current
58 video a user is watching. In their research, Roth et al. focused predominantly
59 on recommendations based on content. They found that such recommendations
60 could quickly lead to a decrease in information diversity (read: filter bubbles)

61 and that this decrease happened sooner for videos with a lot of views; the
62 more views a video had, the less diverse its related recommendations. They
63 speculate that this can be explained by the fact that YouTube tends to store
64 more information about videos with a high view count, allowing the algorithm to
65 give better recommendations for such videos. They also predict that, whenever
66 the algorithm has more information about a user to its disposal, it can combine
67 said information with the information it has about a certain video, which could
68 lead to an even stronger limitation of recommendations. According to Ledwich
69 and Zaitsev (2019), a user's viewing behavior is responsible for approximately
70 70% of their recommendations; this behavior could therefore play a big role in
71 the creation of filter bubbles on YouTube.

2.3 Conspiracy content on YouTube

72 YouTube has limited rules with regards to the spread of conspiracy videos
73 (YouTube, 2021). As long as the content does not directly incite violence or
74 endangers the public health (e.g. misinformation about the COVID-19 virus),
75 objectively incorrect ideas are allowed to be shared on YouTube. As a result,
76 YouTube is a home to multiple conspiracy communities. Conspiracy theories
77 such as 'the earth is flat and the government is hiding it from us', 'the world
78 will end soon and only followers of this specific religion will be spared', and 'the
79 world is ruled by cannibalistic, satanic pedophiles' (better known as QAnon)
80 gather millions of views on the platform (Paolillo, 2018; Miller, 2021). Though
81 such videos could be considered harmful to society, they are not suppressed by
82 YouTube. Whenever a user shows interest in this type of content, they will
83 be recommended similar videos, even when YouTube is aware of their harmful
84 nature (Ledwich and Zaitsev, 2019; Maack, 2019).

2.4 The YouTube algorithm

85 The YouTube algorithm tries to recommend videos based on the expected watch
86 time they will generate, rather than the probability of a user clicking on them
87 (Covington et al., 2016). This decision was made in order to decrease the
88 likelihood of misleading videos (also known as *clickbait*) being recommended.
89 However, gathering feedback about videos through their watch time can cause
90 a lot of noise, making it difficult to measure user satisfaction. As it turns
91 out, even when a users enjoy a certain video, they are unlikely to watch it
92 completely. On average, users watch around 50-60% of a video before they
93 switch it off (Park et al., 2016). Though, videos that are well-structured, or
94 especially interesting, can improve this percentage up till 70-80%, where nearly
95 half of the viewers actually finish the video in its entirety (Lang, 2018). After a
96 video has been watched, there is a 41.6% chance that the user decides to watch
97 a recommended video. Which recommendation the user will choose, follows a
98 Zipf-distribution ($\alpha = 0.78$) with regards to the position of the video in the list
99 of recommendations (Zhou et al., 2010).

100 All in all, previous research has found that YouTube's algorithm is
101 sensitive to filter bubbles and that it has a tendency to recommend conspiracy
102 content. It is also speculated that the algorithm makes decisions based on the
103 user's viewing behavior, which it combines with the content of videos. In order to
104 keep the user on the website as long as possible, which is profitable for YouTube,
105 the algorithm prefers recommending videos that it suspects the user will watch
106 for a longer period of time, even when they may contain harmful content. Based
107 on this information, further research can be done on the origination of filter
108 bubbles and the spread of conspiracy content on YouTube. For example, little is
109 known about how quickly a user's recommendations adapt to a user's behavior,
110 even though this is a critical aspect when it comes to the creation of so-called
111 *rabbit holes*. Furthermore, no research has been done into the way different
112 types of videos (recommendations in different locations, random videos, etc.)
113 influence YouTube's algorithm. For example, whenever a user primarily watches
114 recommended videos, the algorithm could see this as implicit positive feedback,
115 which could cause a snowball-effect.

3 Methodology

3.1 Watching conspiracy videos

116 In order to determine how different watch strategies affect the YouTube algo-
117 rithm, a python script was created to automatically log into a Google account
118 and proceed to watch YouTube videos. The script was made using Selenium
119 WebDriver: a suite of tools used for browser automation.

3.1.1 Google login

120 Due to Google's strict policy regarding automation within their ecosystem, many
121 obstacles are put into place to prevent users from logging into a Google account
122 using automated software such as a selenium script. To circumvent this restric-
123 tion, two steps had to be taken. Firstly, the selenium WebDriver had to be
124 accompanied by the selenium-stealth package, which removes metadata about
125 the current browser, so that it is less obvious that a WebDriver is being used.
126 Additionally, because this metadata was removed, the Google login service was
127 unable to check what browser the client was using. This results in a warning
128 to the user that their current browser may be insecure, which prohibits them
129 from logging in. To avoid this warning, the Google account needs to have been
130 created within a WebDriver, such as Google's ChromeDriver or Mozilla's Gecko-
131 Driver. Therefore, all twenty accounts were manually created in ChromeDriver.
132 Since Google accounts require a phone number verification upon creation, six
133 free (prepaid) SIM cards were ordered from various providers in order to create
134 the accounts. Each SIM card could create two to three accounts before it was
135 blocked due to being used too many times.

3.1.2 The watch strategies

136 After all accounts had been created, they were subdivided into four distinct
137 watch strategies, making for a total of five accounts per strategy. The first
138 watch strategy is the simplest one. The bots following it will watch random,
139 non-conspiracy videos from a dataset. This watch strategy is used as the base-
140 line to compare the other three strategies to. The second strategy is similar
141 to the first: the adhering bots watch random conspiracy videos from a dataset.
142 This strategy was not added to study the origination of filter bubbles, since it is
143 unlikely that a filter bubble will come from this strategy. Considering the video
144 choices will be all over the place, it will be difficult for the algorithm to deter-
145 mine the specific interest of the user. However, comparing how the algorithm
146 responds to conspiracy content in general as opposed to regular content might
147 still yield interesting results. The second-to-last strategy starts off in the same
148 way: it chooses a random conspiracy video from a dataset to watch. However,
149 after having watched the initial video, it will start watching the recommended
150 videos displayed next to the current video. These recommendations consist of
151 a combination of recommendations based on the content of the current video
152 and the personalized recommendations of the user. By using this strategy, bots
153 are likely to go *down the rabbit hole* and eventually end up in a filter bubble.
154 Finally, the last strategy will be similar to the previous one, though with one
155 alteration: rather than choosing a recommended video from the list of recom-
156 mendations next to the current video, it will choose a recommended video from
157 the YouTube homepage of the account. Compared to the third strategy, this
158 will lead to the user watching more personalized recommendations rather than
159 content-based recommendations, possibly speeding up the creation and/or in-
160 creasing the strength of the filter bubble. Each strategy will be executed by
161 five different accounts in order to decrease the probability of a random streak of
162 videos altering the result. The individual accounts will watch a total of fifteen
163 videos as described by their watch strategy, for a total of three hundred videos
164 watched by the script.

3.1.3 Running the bots

165 After the accounts were logged in, they started watching YouTube videos ac-
166 cording to their watch strategy. However, some restrictions were put into place
167 to make sure the bots did not take too long (considering three hundred videos
168 had to be watched in total, some limitations had to apply). For example, the
169 bots were not allowed to watch videos over an hour long, nor were they allowed
170 to watch live streams, as those could theoretically go on infinitely. Additionally,
171 the random videos at the start of the third and fourth strategy were first manu-
172 ally inspected to make sure the bots would not start the experiment by watching
173 a falsely flagged conspiracy video. Considering the way in which the dataset
174 was created, it is possible that some videos that are flagged as conspiracy videos
175 are, in reality, normal videos. This could happen whenever a conspiracy chan-
176 nel uploads a regular video for once (e.g. a holiday video, promotion of some

177 product, etc.). These *false positives* are far and few between, however, having
178 one of them be selected as the first video for strategy three or four had to be
179 prevented, as that could greatly alter the final results. With these restrictions
180 in mind, the following script was created and run for all twenty bots, keeping
181 track of the videos they watched and the homepage recommendations they had
182 after each video:

Algorithm 1: Watch YouTube videos according to a watch strategy

Data: User information and a video dataset

Result: The watched videos and homepage recommendations of the user

```
1 initialize WebDriver;
2 log into Google account;
3 for twenty videos do
4   if there is a recommendation to be watched then
5     go to the link;
6   else
7     pick a random video to watch based on usertype;
8     determine how long it will get watched;
9     go to the link;
10  get video metadata and store for overview of watched videos;
11  watch video for given amount of time;
12  if usertype == 3 then
13    pick recommendation next to current video to watch next;
14    determine watch time for found recommendation;
15  go to YouTube homepage;
16  store current recommendations for overview;
17  if usertype == 4 then
18    pick homepage recommendation to watch next;
19    determine watch time for found recommendation;
20 return watched videos and homepage recommendations;
```

Running the script for all twenty bots resulted in two different datasets: the first containing the videos watched by the bots, including their view count, likes/dislikes, and upload date; and one containing the homepage recommendations for all bots, after each number of videos watched. To determine the influence of the watch strategies on the algorithm, the recommendations were labeled as being either conspiracy or non-conspiracy videos by the classifier. By then grouping all recommendations by their watch strategy and the number of videos watched before them (i.e. the recommendations after the third video watched by all bots with strategy one), it was possible to calculate aggregates

about general statistics of the recommendations, such as view count and video duration, and the percentage of conspiracy videos present amongst them. This lead to four groups with fifteen entries of different statistics (one for each video watched). In order to find out whether any of the differences between the four groups were significant at any point, a number of ANOVAs were performed.

3.2 Leaving the filter bubble

3.3 Machine learning

3.3.1 Data gathering

183 To answer the research question, it is necessary to determine which YouTube
184 videos can be considered conspiracy videos. Considering the large amount of
185 videos getting recommended, determining each video manually is simply not
186 possible. There are two possible ways to solve this problem. Firstly, there is a
187 dataset which contains nearly 7000 YouTube channels that have been manually
188 labeled based on their political view - almost 3000 of which were labeled as
189 conspiracy channels (Ledwich and Zaitsev, 2019); whenever a video is made
190 by one such channel, it can be considered a conspiracy video. However, due
191 to the enormous amount of existing YouTube channels, the odds of a video
192 being uploaded by a channel that is not present in this dataset are very large.
193 For those videos, a supervised machine learning classifier was used. To optimize
194 performance, five different classifiers have been trained and compared: k-nearest
195 neighbors, support-vector machine, neural network, logistic regression, and ridge
196 regression.

197 In order to train these machine learning algorithms, a training dataset
198 was created. To get a labeled dataset of conspiracy and non-conspiracy videos,
199 use was made of the aforementioned channel dataset made by Ledwich and
200 Zaitsev (2019). For each channel in that dataset, the title, description, and
201 transcript of the ten most recently uploaded videos were downloaded using
202 YouTube's API. Videos uploaded by a conspiracy channel were then labeled
203 as conspiracy videos, and videos uploaded by a channel from a different category
204 were labeled as normal videos. Additionally, the channel description and
205 channel keywords (which are used for targeted advertising on YouTube) were
206 added to each video. The final dataset contains 65.683 unique YouTube videos,
207 22.156 of which are considered as conspiracy videos.

3.3.2 Data cleaning

208 However, this dataset was not yet suitable for machine learning, as the data
209 was still messy. Therefore, multiple steps were taken in order to clean the data.
210 Firstly, the two classes (conspiracy and non-conspiracy) were balanced, so that
211 the classifier would not develop a bias for non-conspiracy videos. Rather than
212 opting for balancing the two classes through the use of class-weights (a technique
213 where weights are attributed to classes, thereby telling the classifier that getting
214 a prediction correct for a certain, underrepresented class is more important), the

choice was made to under-sample the data in order to equalize both classes (both containing 22.156 videos, for a total of 44.312 videos) (Lemaître et al., 2017; Sun et al., 2006). As there was plenty of data in the dataset, under-sampling was more convenient than implementing class-weights. After both classes had been balanced, the text for each video had to be translated into English. Since the original dataset by Ledwich and Zaitsev (2019) also contained channels by non-English speakers, these videos had to be automatically translated. Then, a few common cleaning methods were applied: all text was converted to lowercase, after which special characters, such as emojis were removed, whereafter stop words were removed and all words were stemmed using the porter stemmer (Karaa, 2013). Finally, each video was TF-IDF vectorized to allow the classifiers to function.

3.3.3 Performance optimization

After splitting the dataset into a training, test, and validation set, the hyperparameters of each algorithm were tuned to get the optimal performance (Feurer and Hutter, 2019). Performance was measured using four distinct metrics: the accuracy, which shows the share of correct predictions; the recall, which shows what fraction of truly positive samples were correctly labeled as such; the precision, which shows what part of the positive predictions were correct; and the F1-score, which is the harmonic mean of the recall and precision (Sokolova and Lapalme, 2009). For each classifier, different configurations of hyperparameters (such as the kernel and the penalty-parameter) were systemically tested - each possible combination was tried. The classifiers were trained on the training set and the optimal hyperparameters were determined based on the performance of the classifiers on the validation set. By saving these performance measures for every configuration, for every classifier, the optimal configuration for every classifier could be determined. Lastly, the classifiers were equipped with their optimal hyperparameters and then tested for the final time on the test set. By comparing the performance of every optimally configured classifier on the test set, the best-performing classifier could be chosen (Reitermanova, 2010).

Additionally, the added value of using a machine learning ensemble was measured. By having each classifier make a prediction for all videos in the dataset, a new dataset was created, wherein the features were the predictions made by the different classifiers. By using all possible combinations of classifiers, and then having different neural networks use those features as input, a machine learning ensemble was created. This ensemble was then optimized in a similar way to the classifiers individually.

Classifier performance




Figure 1: Metrics for each classifier with optimized hyperparameters

4 Results

4.1 The recommended content

4.2 Leaving the filter bubble

4.3 Machine learning

251 The hyperparameter tuning lead to impressive scores for all classifiers. When
 252 making prediction for the test set, the best-performing classifier is the support-
 253 vector machine making use of the Radial Basis Function (RBF) kernel and a
 254 penalty parameter (C-value) of 10. In second place, there is a two-way tie for
 255 accuracy between ridge regression with a sparse-cg solver and penalty (alpha)
 256 value of 0.1, and the neural network using the identity activation function with
 257 10 hidden layers of 10 neurons. However, ridge regression has a slightly better
 258 F1-score, though this difference is neglectable (0.9193 as opposed to 0.9184).
 259 The neural network and ridge regression are closely followed by logistic regres-
 260 sion with an L2 penalty, a penalty (C) value of 20 and a newton-cg solver.
 261 The worst-performing classifier is also the simplest of the bunch: the k-nearest
 262 neighbors classifier (K=1). Although its performance is still formidable, it does
 263 substantially worse than the others. An overview of all metrics for each clas-
 264 sifier can be seen in figure 1. The ten best-performing configurations for each

265 classifier can be found in appendix A.

266 Noteworthy is the fact that the optimal ensemble actually outperforms
267 the support-vector machine by a slight margin. This ensemble, consisting of the
268 SVM, the neural network, and surprisingly, the k-nearest neighbor classifiers,
269 gets slightly higher scores than the runner-up across the board. The ensemble
270 had a 16-way tie for best-performing parameters, all of which contained at least
271 the SVM, neural network, and k-NN classifiers.

272 Though the ensemble outperforms the other classifiers, it has a signif-
273 icant drawback: its training time is significantly larger than that of the individ-
274 ual classifiers. Support-vector machines are infamous for their slowness when
275 there is a lot of training data, and neural networks can require a lot of training
276 time whenever the number of neurons gets large (Burges and Schölkopf, 1997;
277 Kamarathi and Pittner, 1999). Requiring both algorithms to run will therefore
278 require a lot of additional training time. Considering the marginal performance
279 increase, the cost outweighs the benefit. As a result, when taking everything into
280 account, the support-vector machine is the best classifier for labeling conspiracy
281 videos on YouTube.

5 Discussion

282 To do.



6 Planning

Table 1: Planning

Week	Handelingen	Afgehandeld
28/03-03/04	Hyperparameters optimaliseren	Ja
03/04-10/04	Classifier-ensemble optimaliseren	Ja
11/04-17/04	Google accounts maken, deelvraag 3 maken	Ja
18/04-24/04	Inleiding uitbreiden	Ja
25/04-01/05	Uitvoering experiment, labelen met classifier	
02/05-08/05	Deelvraag 1 schrijven, beginnen deelvraag 2	
09/05-15/05	Deelvraag 2 afschrijven	
16/05-22/05	Beginnen met discussie schrijven	
23/05-29/05	Discussie afschrijven	
30/05-05/06	Abstract schrijven, tekst proof-readen	
06/06-12/06	Laatste aanpassingen en verbeteringen	
13/06-19/06	Inleveren scriptie en verdediging	



References

- 283 Burges, C. J. and Schölkopf, B. (1997). Improving the accuracy and speed
284 of support vector machines. *Advances in neural information processing*
285 *systems*, pages 375–381.
- 286 Cooper, P. (2020). How does the youtube algorithm work? a guide to getting
287 more views.
- 288 Covington, P., Adams, J., and Sargin, E. (2016). Deep neural networks for
289 youtube recommendations. In *Proceedings of the 10th ACM conference on*
290 *recommender systems*, pages 191–198.
- 291 Donzelli, G., Palomba, G., Federigi, I., Aquino, F., Cioni, L., Verani, M., Car-
292 ducci, A., and Lopalco, P. (2018). Misinformation on vaccination: A quan-
293 titative analysis of youtube videos. *Human vaccines & immunotherapeutics*,
294 14(7):1654–1659.
- 295 Feurer, M. and Hutter, F. (2019). Hyperparameter optimization. In *Automated*
296 *Machine Learning*, pages 3–33. Springer, Cham.
- 297 Kamarathi, S. V. and Pittner, S. (1999). Accelerating neural network training
298 using weight extrapolations. *Neural networks*, 12(9):1285–1299.
- 299 Karaa, W. B. A. (2013). A new stemmer to improve information retrieval.
300 *International Journal of Network Security & Its Applications*, 5(4):143.
- 301 Lang, P. (2018). Youtube average view duration - the 50% rule.
- 302 Ledwich, M. and Zaitsev, A. (2019). Algorithmic extremism: Examining
303 youtube’s rabbit hole of radicalization. *arXiv preprint arXiv:1912.11211*.
- 304 Lemaître, G., Nogueira, F., and Aridas, C. K. (2017). Imbalanced-learn: A
305 python toolbox to tackle the curse of imbalanced datasets in machine learn-
306 ing. *The Journal of Machine Learning Research*, 18(1):559–563.
- 307 Lubach, A. (2020). De online fabeltjesfuik — zondag met lubach (s12).
- 308 Maack, M. M. (2019). ‘youtube recommendations are toxic,’ says dev who
309 worked on the algorithm.
- 310 Miller, D. T. (2021). Characterizing qanon: Analysis of youtube comments
311 presents new conclusions about a popular conservative conspiracy. *First*
312 *Monday*.
- 313 Neufeld, D. (2021). The 50 most visited websites in the world.
- 314 Paolillo, J. C. (2018). The flat earth phenomenon on youtube. *First Monday*.
- 315 Pariser, E. (2011). *The filter bubble: What the Internet is hiding from you*.
316 Penguin UK.



- 317 Park, M., Naaman, M., and Berger, J. (2016). A data-driven study of view
318 duration on youtube. In *Proceedings of the International AAAI Conference*
319 *on Web and Social Media*, volume 10.
- 320 Reitermanova, Z. (2010). Data splitting. In *WDS*, volume 10, pages 31–36.
- 321 Rosenbaum, L. (2021). Escaping catch-22—overcoming covid vaccine hesitancy.
- 322 Roth, C., Mazières, A., and Menezes, T. (2020). Tubes and bubbles topological
323 confinement of youtube recommendations. *PloS one*, 15(4):e0231703.
- 324 Sokolova, M. and Lapalme, G. (2009). A systematic analysis of performance
325 measures for classification tasks. *Information processing & management*,
326 45(4):427–437.
- 327 Sun, Y., Kamel, M. S., and Wang, Y. (2006). Boosting for learning multiple
328 classes with imbalanced class distribution. In *Sixth international conference*
329 *on data mining (ICDM'06)*, pages 592–602. IEEE.
- 330 YouTube (2021). Youtube community guidelines & policies - how youtube works.
- 331 Zhou, R., Khemmarat, S., and Gao, L. (2010). The impact of youtube rec-
332 ommendation system on video views. In *Proceedings of the 10th ACM*
333 *SIGCOMM conference on Internet measurement*, pages 404–410.



Appendices

A Hyperparameter tuning

Ensemble	Activation	Layers	Neurons	Accuracy	Precision	Recall	F1
svm, nn, knn	logistic	1	20	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	relu	1	20	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	identity	1	1	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	identity	10	1	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	logistic	1	10	0.939318	0.948923	0.931204	0.93998
ridge, svm, nn, knn	tanh	10	1	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	tanh	1	10	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	tanh	1	20	0.939318	0.948923	0.931204	0.93998
svm, nn, logr, knn	identity	1	1	0.939318	0.948923	0.931204	0.93998
svm, nn, logr, knn	identity	1	10	0.939318	0.948923	0.931204	0.93998

Table 2: Ensemble.

Kernel	C	Accuracy	Precision	Recall	F1
rbf	10.0	0.936309	0.945473	0.928747	0.937035
rbf	100.0	0.935557	0.942289	0.930713	0.936465
rbf	1.0	0.925276	0.930163	0.922850	0.926492
poly	10.0	0.916499	0.946017	0.886978	0.915547
poly	100.0	0.915246	0.944940	0.885504	0.914257
linear	1.0	0.913741	0.917944	0.912531	0.915229
poly	1.0	0.909729	0.935065	0.884521	0.909091
linear	10.0	0.904965	0.907882	0.905651	0.906765
linear	100.0	0.898195	0.905830	0.893366	0.899555
rbf	0.1	0.878887	0.878906	0.884521	0.881705

Table 3: Support-vector machine.

Activation	Layers	Neurons	Accuracy	Precision	Recall	F1
identity	10	10	0.923019	0.935484	0.912039	0.923613
identity	25	10	0.921013	0.933031	0.910565	0.921661
relu	10	10	0.919007	0.917561	0.924324	0.920930
identity	10	20	0.916750	0.906056	0.933661	0.919652
relu	10	20	0.915998	0.912221	0.924324	0.918233
tanh	10	10	0.915747	0.914592	0.920885	0.917728
relu	1	1	0.915747	0.931876	0.900737	0.916042
tanh	25	20	0.915496	0.919052	0.914988	0.917016
tanh	10	20	0.914744	0.920178	0.912039	0.916091
logistic	1	1	0.913741	0.925516	0.903686	0.914470

Table 4: Neural network.

Solver	Alpha	Accuracy	Precision	Recall	F1
auto	0.1	0.918506	0.919118	0.921376	0.920245
sparse_cg	0.1	0.918506	0.919118	0.921376	0.920245
sag	0.1	0.918255	0.919902	0.919902	0.919902
auto	1.0	0.917252	0.923497	0.913514	0.918478
sparse_cg	1.0	0.917252	0.923497	0.913514	0.918478
sag	1.0	0.917252	0.923497	0.913514	0.918478
sag	10.0	0.878385	0.893002	0.865356	0.878962
auto	10.0	0.878134	0.892549	0.865356	0.878743
sparse_cg	10.0	0.878134	0.892549	0.865356	0.878743
auto	100.0	0.812437	0.854545	0.762162	0.805714

Table 5: Ridge regression.

Penalty	C	Solver	Accuracy	Precision	Recall	F1
l2	20	newton-cg	0.918506	0.924107	0.915479	0.919773
l2	20	saga	0.918506	0.924107	0.915479	0.919773
l2	20	sag	0.918506	0.924107	0.915479	0.919773
l2	10	sag	0.916249	0.920831	0.914496	0.917653
l2	10	newton-cg	0.916249	0.920831	0.914496	0.917653
l2	10	saga	0.916249	0.920831	0.914496	0.917653
l2	10	lbfgs	0.915747	0.919506	0.914988	0.917241
l2	20	lbfgs	0.914744	0.921432	0.910565	0.915966
none	1	sag	0.913992	0.923848	0.906143	0.914909
none	10	saga	0.913240	0.922461	0.906143	0.914229

Table 6: Logistic regression.



K	Accuracy	Precision	Recall	F1
1	0.889669	0.888456	0.896314	0.892368
3	0.888415	0.882212	0.901720	0.891859
4	0.879137	0.908899	0.848157	0.877478
5	0.873370	0.858482	0.900246	0.878868
6	0.873119	0.882441	0.866830	0.874566
2	0.872618	0.935043	0.806388	0.865963
7	0.868355	0.848891	0.902703	0.874970
8	0.867603	0.867382	0.874201	0.870778
9	0.861585	0.835672	0.907125	0.869934
10	0.859579	0.854397	0.873710	0.863946

Table 7: K-nearest neighbors.