



UNIVERSITY OF AMSTERDAM

## BACHELOR THESIS

---

# The impact of conspiracy videos on the YouTube algorithm

---

April 28, 2021

*Student:*

Roan Schellingerhout

*Supervisor:*

Dr. Maarten Marx

### Abstract

Vivamus eu tellus sed tellus consequat suscipit. Nam orci orci, malesuada id, gravida nec, ultricies vitae, erat. Donec risus turpis, luctus sit amet, interdum quis, porta sed, ipsum. Suspendisse condimentum, tortor at egestas posuere, neque metus tempor orci, et tincidunt urna nunc a purus. Sed facilisis blandit tellus. Nunc risus sem, suscipit nec, eleifend quis, cursus quis, libero. Curabitur et dolor. Sed vitae sem. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Maecenas ante. Duis ullamcorper enim. Donec tristique enim eu leo. Nullam molestie elit eu dolor. Nullam bibendum, turpis vitae tristique gravida, quam sapien tempor lectus, quis pretium tellus purus ac quam. Nulla facilisi.

## 1 Introduction

1 YouTube attracts an average of 34.6 billion page views per month, making it  
2 the world's largest video-sharing website and the second largest website on the  
3 entire internet (Neufeld, 2021). The overwhelming majority of those page views  
4 come from users watching videos, 70% of which are recommended to users by  
5 YouTube's algorithm (Cooper, 2020). All types of content get produced and  
6 consumed on the website. However, conspiracy content has been booming on  
7 YouTube (Donzelli et al., 2018). Alt-right (or far-right) and conspiracy chan-  
8 nels are starting to grow their audiences, which could have many negative con-  
9 sequences for society at large. For example, the number of people who are  
10 distrustful of science is increasing, a development in which conspiracy content  
11 on YouTube plays a role. Whenever this increased distrust relates to impor-  
12 tant topics, such as believing in the efficacy of vaccines, it can create genuine

dangers to the public. As it turns out, more than half of the American population has doubts about - or is definitely against - taking the COVID-19 vaccine (Rosenbaum, 2021).

To better understand how YouTube's algorithm (and recommender algorithms in general) allow(s) conspiracy content to thrive, this research will investigate how quickly the algorithm develops a preference for conspiracy video; in other words: how many videos a user needs to watch before they get sent *down the rabbit hole*.

This research takes inspiration from the Dutch television program *Zondag met Lubach*, where a similar idea was executed on a smaller scale (Lubach, 2020). Lubach's experiment consisted of creating a new YouTube account on a never-used laptop, after which the account was used to watch a few (recommended) conspiracy videos, to see how this would affect the YouTube homepage of the account. After a mere three videos, the homepage of the account was filled to the brim with conspiracy content, mostly having to do with the coronavirus. The extremely interesting results that came from this micro-experiment formed the motivation to research this phenomenon more extensively.

## 1.1 Research question

To research this subject, the following research question has been formulated: *What is the impact of different watch strategies on the number of conspiracy videos that have to be watched until a user's YouTube-recommendations start preferring conspiracy content?* In this scenario, 'preferring' will be defined as the situation in which the amount of conspiracy videos present in the recommendations is significantly higher than that of the baseline.

To assist in answering the research question, the following sub-questions will be answered:

- How do different watch strategies on YouTube influence the type of content that is recommended to a user?
- How long does it take for a YouTube recommendations to stop preferring conspiracy videos, once they have started doing so?
- What type of classifier performs the best when it comes to labeling conspiracy videos on YouTube?

## 2 Theoretical Framework

### 2.1 Filter bubbles on social media

Whenever the user of a website finds themselves in their own information universe, in which the content and recommendations play into the user's preexisting opinions and beliefs, they are in a filter bubble (Pariser, 2011). Users are by themselves in such bubbles and each bubble is unique. Different bubbles can

49 have overlap, but each bubble is precisely tuned to an individual. In traditional  
50 media, a user makes a conscious *choice* what types of opinions they want to  
51 hear, for example by choosing to watch a broadcaster with a specific political  
52 opinion. Online this decision is implicit: based on the user's behavior, their  
53 content is filtered automatically by an algorithm, without explicit consent.

## 2.2 Filter bubbles on YouTube

54 Previous research has found that YouTube's recommendation algorithm runs  
55 the risk of creating filter bubbles. Roth et al. (2020) came to this conclusion  
56 after they analysed YouTube recommendations based on content. YouTube has  
57 two distinct types of recommendations: recommendations based on the user's  
58 viewing behavior and recommendations based on the content of the current  
59 video a user is watching. In their research, Roth et al. focused predominantly  
60 on recommendations based on content. They found that such recommendations  
61 could quickly lead to a decrease in information diversity (read: filter bubbles)  
62 and that this decrease happened sooner for videos with a lot of views; the  
63 more views a video had, the less diverse its related recommendations. They  
64 speculate that this can be explained by the fact that YouTube tends to store  
65 more information about videos with a high view count, allowing the algorithm to  
66 give better recommendations for such videos. They also predict that, whenever  
67 the algorithm has more information about a user to its disposal, it can combine  
68 said information with the information it has about a certain video, which could  
69 lead to an even stronger limitation of recommendations. According to Ledwich  
70 and Zaitsev (2019), a user's viewing behavior is responsible for approximately  
71 70% of their recommendations; this behavior could therefore play a big role in  
72 the creation of filter bubbles on YouTube.

## 2.3 Conspiracy content on YouTube

73 YouTube has limited rules with regards to the spread of conspiracy videos  
74 (YouTube, 2021). As long as the content does not directly incite violence or  
75 endangers the public health (e.g. misinformation about the COVID-19 virus),  
76 objectively incorrect ideas are allowed to be shared on YouTube. As a result,  
77 YouTube is a home to multiple conspiracy communities. Conspiracy theories  
78 such as 'the earth is flat and the government is hiding it from us', 'the world  
79 will end soon and only followers of this specific religion will be spared', and 'the  
80 world is ruled by cannibalistic, satanic pedophiles' (better known as QAnon)  
81 gather millions of views on the platform (Paolillo, 2018; Miller, 2021). Though  
82 such videos could be considered harmful to society, they are not suppressed by  
83 YouTube. Whenever a user shows interest in this type of content, they will  
84 be recommended similar videos, even when YouTube is aware of their harmful  
85 nature (Ledwich and Zaitsev, 2019; Maack, 2019).

## 2.4 The YouTube algorithm

86 The YouTube algorithm tries to recommend videos based on the expected watch  
87 time they will generate, rather than the probability of a user clicking on them  
88 (Covington et al., 2016). This decision was made in order to decrease the  
89 likelihood of misleading videos (also known as *clickbait*) being recommended.  
90 However, gathering feedback about videos through their watch time can cause  
91 a lot of noise, making it difficult to measure user satisfaction. As it turns  
92 out, even when a users enjoy a certain video, they are unlikely to watch it  
93 completely. On average, users watch around 50-60% of a video before they  
94 switch it off (Park et al., 2016). Though, videos that are well-structured, or  
95 especially interesting, can improve this percentage up till 70-80%, where nearly  
96 half of the viewers actually finish the video in its entirety (Lang, 2018). After a  
97 video has been watched, there is a 41.6% chance that the user decides to watch  
98 a recommended video. Which recommendation the user will choose, follows a  
99 Zipf-distribution ( $\alpha = 0.78$ ) with regards to the position of the video in the list  
100 of recommendations (Zhou et al., 2010).

101 All in all, previous research has found that YouTube's algorithm is  
102 sensitive to filter bubbles and that it has a tendency to recommend conspiracy  
103 content. It is also speculated that the algorithm makes decisions based on the  
104 user's viewing behavior, which it combines with the content of videos. In order to  
105 keep the user on the website as long as possible, which is profitable for YouTube,  
106 the algorithm prefers recommending videos that it suspects the user will watch  
107 for a longer period of time, even when they may contain harmful content. Based  
108 on this information, further research can be done on the origination of filter  
109 bubbles and the spread of conspiracy content on YouTube. For example, little is  
110 known about how quickly a user's recommendations adapt to a user's behavior,  
111 even though this is a critical aspect when it comes to the creation of so-called  
112 *rabbit holes*. Furthermore, no research has been done into the way different  
113 types of videos (recommendations in different locations, random videos, etc.)  
114 influence YouTube's algorithm. For example, whenever a user primarily watches  
115 recommended videos, the algorithm could see this as implicit positive feedback,  
116 which could cause a snowball-effect.

## 3 Methodology

### 3.1 Watching conspiracy videos

117 In order to determine how different watch strategies affect the YouTube algo-  
118 rithm, a python script was created to automatically log into a Google account  
119 and proceed to watch YouTube videos. The script was made using Selenium  
120 WebDriver: a suite of tools used for browser automation.

121 In the following sections, each aspect of the initial experiment will be  
122 explained. Firstly, an explanation will be given about how to log into Google  
123 accounts using a python bot. Then, the watch strategies as mentioned in the  
124 research question will be defined, followed by a description of their video watch-

125 ing behavior. Afterwards, some restrictions about the videos being watched will  
126 be mentioned. Additionally the actual script allowing the bots to be run will be  
127 described. Lastly, the precise form of the output of the script will be explained.

### 3.1.1 Google login

128 Due to Google's strict policy regarding automation within their ecosystem, many  
129 obstacles are put into place to prevent users from logging into a Google account  
130 using automated software such as a selenium script. To circumvent this restric-  
131 tion, two steps had to be taken. Firstly, the selenium WebDriver had to be  
132 accompanied by the selenium-stealth package, which removes metadata about  
133 the current browser, so that it is less obvious that a WebDriver is being used.  
134 Additionally, because this metadata was removed, the Google login service was  
135 unable to check what browser the client was using. This results in a warning  
136 to the user that their current browser may be insecure, which prohibits them  
137 from logging in. To avoid this warning, the Google account needs to have been  
138 created within a WebDriver, such as Google's ChromeDriver or Mozilla's Gecko-  
139 Driver. Therefore, all twenty accounts were manually created in ChromeDriver.  
140 Since Google accounts require a phone number verification upon creation, six  
141 free (prepaid) SIM cards were ordered from various providers in order to create  
142 the accounts. Each SIM card could create two to three accounts before it was  
143 blocked due to being used too many times.

### 3.1.2 The watch strategies

144 After all accounts had been created, they were subdivided into four distinct  
145 watch strategies, making for a total of five accounts per strategy.

146 **Random videos (baseline)** The first watch strategy is the simplest one. The  
147 bots following it will watch random, non-conspiracy videos from a dataset.  
148 This watch strategy is used as the baseline to compare the other three  
149 strategies to.

150 **Random conspiracies** The second strategy is similar to the first: the adhering  
151 bots watch random conspiracy videos from a dataset. This strategy was  
152 not added to study the origination of filter bubbles, since it is unlikely  
153 that a filter bubble will come from this strategy. Considering the video  
154 choices will be all over the place, it will be difficult for the algorithm to  
155 determine the specific interest of the user. However, comparing how the  
156 algorithm responds to conspiracy content in general as opposed to regular  
157 content might still yield interesting results.

158 **Video recommendations** The second-to-last strategy starts off in the same  
159 way: it chooses a random conspiracy video from a dataset to watch. How-  
160 ever, after having watched the initial video, it will start watching the  
161 recommended videos displayed next to the current video. These recom-  
162 mendations consist of a combination of recommendations based on the

content of the current video and the personalized recommendations of the user. By using this strategy, bots are likely to go *down the rabbit hole* and eventually end up in a filter bubble.

**Homepage recommendations** Finally, the last strategy will be similar to the previous one, though with one alteration: rather than choosing a recommended video from the list of recommendations next to the current video, it will choose a recommended video from the YouTube homepage of the account. Compared to the third strategy, this will lead to the user watching more personalized recommendations rather than content-based recommendations, possibly speeding up the creation and/or increasing the strength of the filter bubble.

Each strategy will be executed by five different accounts in order to decrease the probability of a random streak of videos altering the result. The individual accounts will watch a total of fifteen videos as described by their watch strategy, for a total of three hundred videos watched by the script.

Additionally, to simulate real-world user behavior, the average watch time for the videos will be normally distributed with a mean of 55% and a standard deviation of 25% (Park et al., 2016; Lang, 2018). The watch time will not be able to exceed a value of 100, as a video cannot be watched for more than 100%. In the same vein, the clicking behavior of users will be simulated as accurately as possible. The probability of a user clicking on a video at position  $k$  within a given list of recommendations (its click-through rate:  $CTR$ ), will be determined using the following formula:

$$CTR(k; N, \alpha) = \frac{1/k^\alpha}{\sum_{n=1}^N (1/n^\alpha)} \quad (1)$$

Wherein  $N$  is the total number of recommendations and  $\alpha$  is the distribution's exponent value ( $\alpha = 0.78$ ) (Zhou et al., 2010). Using this formula, when considering the first twenty recommendations, the first recommendation will have a click-through rate of approximately 20.6%, after which the CTR quickly decreases, until a probability of 1.9% at the twentieth recommendation.

### 3.1.3 Running the bots

After the accounts were logged in, they started watching YouTube videos according to their watch strategy. However, some restrictions were put into place to make sure the bots did not take too long (considering three hundred videos had to be watched in total, some limitations had to apply). For example, the bots were not allowed to watch videos over an hour long, nor were they allowed to watch live streams, as those could theoretically go on infinitely. Additionally, the random videos at the start of the third and fourth strategy were first manually inspected to make sure the bots would not start the experiment by watching a falsely flagged conspiracy video. Considering the way in which the dataset was created, it is possible that some videos that are flagged as conspiracy

201 videos are, in reality, normal videos. This could happen whenever a conspiracy  
202 channel uploads a regular video for once (e.g. a holiday video, promotion of  
203 some product, etc.). These *false positives* are far and few between, however,  
204 having one of them be selected as the first video for strategy three or four had  
205 to be prevented, as that could have greatly altered the final results. With these  
206 restrictions in mind, the following script was created and run for all twenty bots,  
207 keeping track of the videos they watched and the homepage recommendations  
208 they had after each video:

---

**Algorithm 1:** Watch YouTube videos according to a watch strategy

---

**Data:** User information and a video dataset

**Result:** The watched videos and homepage recommendations of the user

```
1 for twenty bots do
2   initialize WebDriver;
3   log into Google account;
4   for fifteen videos do
5     if there is a recommendation to be watched then
6       go to the link;
7     else
8       pick a random video to watch based on usertype;
9       determine how long it will get watched;
10      go to the link;
11
12    get video metadata and store for overview of watched videos;
13    watch video for given amount of time;
14
15    if usertype == 3 then
16      pick recommendation next to current video to watch next;
17      determine watch time for found recommendation;
18
19    go to YouTube homepage;
20    store current recommendations for overview;
21
22    if usertype == 4 then
23      pick homepage recommendation to watch next;
24      determine watch time for found recommendation;
25
26 return watched videos and homepage recommendations;
```

---

209 Running the script for all twenty bots resulted in two different datasets: the  
210 first containing the videos watched by the bots, including their view count,  
211 likes/dislikes, and upload date; and one containing the homepage recommen-  
212 dations for all bots, after each number of videos watched. To determine the  
213 influence of the watch strategies on the algorithm, the recommendations were

### Example DataFrame

bot	vid_no	video_id	views	likes	dislikes	vid_len	title	description	transcript	channel_desc	keywords	conspiracy
1	5	u6sN6K-4_qo	49778	2913	16	536	All I Want for Christmas (Is a Pardon from the...	It's almost Christmas time. Saxophones by Jon ...	It's almost christmas time and i'm all by myse...	Nick Lutsko is a songwriter, producer, and per...	Nick Lutsko Puppets Etc. ALL SHOOK UP Greezy R...	False
1	5	p212D-WvKGO	71645	3073	38	639	How A Narcissist Reacts When You Behave Badly	How a narcissist reacts when you do something ...	[Music] today I'm going to talk about the diff...	Pop Culture -- Narcissism & Sociopathy -- Phil...	Narcissism Sociopathy Philosophy "Trauma Infor...	True
1	5	7vLbNZs2sFk	25137	132	18	439	FORECAST: Phoenix will stay in the 70s though ...	High pressure over the Pacific is producing a ...	let's take a look at what we're tracking weath...	Get all the news of the day, plus investigativ...	cbs5az 3tv "cbs 5 phoenix" "cbs 5 az" ktvk kph...	False

Figure 1: An example of the experiment's output

labeled as being either conspiracy or non-conspiracy videos by the classifier. By then grouping all recommendations by their watch strategy and the number of videos watched before them (e.g. the recommendations after the third video watched by all bots with strategy one), it was possible to calculate aggregates about general statistics of the recommendations, such as view count and video duration, and the percentage of conspiracy videos present amongst them. This lead to four groups with fifteen entries of different statistics (one for each video watched). In order to find out whether any of the differences between the four groups were significant at any point, a number of ANOVAs were performed.

#### 3.1.4 Outputs of the experiment

To create the initial recommendation dataset, the bots stored the following information about their top twenty recommendations after each video watched:

- the id of the bot that had gotten the recommendations;
- the amount of videos watched before the recommendation was given;
- the URL of the video being recommended;
- the URL of the channel that uploaded the recommendation.

To then allow the classifier to label the videos, and to calculate the aggregates per strategy, additional data was gathered using YouTube's API. For each recommendation, the title, description, transcript, (dis)likes, views, video duration, channel description, and channel keywords were collected. The title, description, transcript, channel description, and channel keywords were collected in order for the classifier to label the videos, as will be further explained in section 3.3. The other features were collected in order to calculate the aggregates used for the aforementioned ANOVAs. An example of the output can be seen in 1.



## 3.2 Leaving the filter bubble

## 3.3 Machine learning

### 3.3.1 Data gathering

237 To answer the research question, it is necessary to determine which YouTube  
238 videos can be considered conspiracy videos. Considering the large amount of  
239 videos getting recommended, determining each video manually is simply not  
240 possible. There are two possible ways to solve this problem. Firstly, there is a  
241 dataset which contains nearly 7000 YouTube channels that have been manually  
242 labeled based on their political view - almost 3000 of which were labeled as  
243 conspiracy channels (Ledwich and Zaitsev, 2019); whenever a video is made  
244 by one such channel, it can be considered a conspiracy video. However, due  
245 to the enormous amount of existing YouTube channels, the odds of a video  
246 being uploaded by a channel that is not present in this dataset are very large.  
247 For those videos, a supervised machine learning classifier was used. To optimize  
248 performance, five different classifiers have been trained and compared: k-nearest  
249 neighbors, support-vector machine, neural network, logistic regression, and ridge  
250 regression.

251 In order to train these machine learning algorithms, a training dataset  
252 was created. To get a labeled dataset of conspiracy and non-conspiracy videos,  
253 use was made of the aforementioned channel dataset made by Ledwich and  
254 Zaitsev (2019). For each channel in that dataset, the title, description, and  
255 transcript of the ten most recently uploaded videos were downloaded using  
256 YouTube's API. Videos uploaded by a conspiracy channel were then labeled  
257 as conspiracy videos, and videos uploaded by a channel from a different cate-  
258 gory were labeled as normal videos. Additionally, the channel description and  
259 channel keywords (which are used for targeted advertising on YouTube) were  
260 added to each video. The final dataset contains 65.683 unique YouTube videos,  
261 22.156 of which are considered as conspiracy videos.

### 3.3.2 Data cleaning

262 However, this dataset was not yet suitable for machine learning, as the data  
263 was still messy. Therefore, multiple steps were taken in order to clean the data.  
264 Firstly, the two classes (conspiracy and non-conspiracy) were balanced, so that  
265 the classifier would not develop a bias for non-conspiracy videos. Rather than  
266 opting for balancing the two classes through the use of class-weights (a technique  
267 where weights are attributed to classes, thereby telling the classifier that getting  
268 a prediction correct for a certain, underrepresented class is more important), the  
269 choice was made to under-sample the data in order to equalize both classes (both  
270 containing 22.156 videos, for a total of 44.312 videos) (Lemaître et al., 2017; Sun  
271 et al., 2006). As there was plenty of data in the dataset, under-sampling was  
272 more convenient than implementing class-weights. After both classes had been  
273 balanced, the text for each video had to be translated into English. Since the  
274 original dataset by Ledwich and Zaitsev (2019) also contained channels by non-

English speakers, these videos had to be automatically translated. Then, a few common cleaning methods were applied: all text was converted to lowercase, after which special characters, such as emojis were removed, whereafter stop words were removed and all words were stemmed using the porter stemmer (Karaa, 2013). Finally, each video was TF-IDF vectorized to allow the classifiers to function.

### 3.3.3 Performance optimization

After splitting the dataset into a training, test, and validation set, the hyperparameters of each algorithm were tuned to get the optimal performance (Feurer and Hutter, 2019). Performance was measured using four distinct metrics: the accuracy, which shows the share of correct predictions; the recall, which shows what fraction of truly positive samples were correctly labeled as such; the precision, which shows what part of the positive predictions were correct; and the F1-score, which is the harmonic mean of the recall and precision (Sokolova and Lapalme, 2009). For each classifier, different configurations of hyperparameters (such as the kernel and the penalty-parameter) were systemically tested - each possible combination was tried. The classifiers were trained on the training set and the optimal hyperparameters were determined based on the performance of the classifiers on the validation set. By saving these performance measures for every configuration, for every classifier, the optimal configuration for every classifier could be determined. Lastly, the classifiers were equipped with their optimal hyperparameters and then tested for the final time on the test set. By comparing the performance of every optimally configured classifier on the test set, the best-performing classifier could be chosen (Reitermanova, 2010).

Additionally, the added value of using a machine learning ensemble was measured. By having each classifier make a prediction for all videos in the dataset, a new dataset was created, wherein the features were the predictions made by the different classifiers. By using all possible combinations of classifiers, and then having different neural networks use those features as input, a machine learning ensemble was created. This ensemble was then optimized in a similar way to the classifiers individually.

## 4 Results

### 4.1 The recommended content

### 4.2 Leaving the filter bubble

### 4.3 Machine learning

The hyperparameter tuning lead to impressive scores for all classifiers. When making predictions for the test set, the best-performing classifier is the support-vector machine making use of the Radial Basis Function (RBF) kernel and a penalty parameter (C-value) of 10. In second place, there is a two-way tie for

## Classifier performance




Figure 2: Metrics for each classifier with optimized hyperparameters

309 accuracy between ridge regression with a sparse-cg solver and penalty (alpha)  
 310 value of 0.1, and the neural network using the identity activation function with  
 311 10 hidden layers of 10 neurons. However, ridge regression has a slightly better  
 312 F1-score, though this difference is neglectable (0.9193 as opposed to 0.9184).  
 313 The neural network and ridge regression are closely followed by logistic regres-  
 314 sion with an L2 penalty, a penalty (C) value of 20 and a newton-cg solver.  
 315 The worst-performing classifier is also the simplest of the bunch: the k-nearest  
 316 neighbors classifier (K=1). Although its performance is still formidable, it does  
 317 substantially worse than the others. An overview of all metrics for each clas-  
 318 sifier can be seen in figure 2. The ten best-performing configurations for each  
 319 classifier can be found in appendix A.

320 Noteworthy is the fact that the optimal ensemble actually outperforms  
 321 the support-vector machine by a slight margin. This ensemble, consisting of the  
 322 SVM, the neural network, and surprisingly, the k-nearest neighbor classifiers,  
 323 gets slightly higher scores than the runner-up across the board. The ensemble  
 324 had a 16-way tie for best-performing parameters, all of which contained at least  
 325 the SVM, neural network, and k-NN classifiers.

326 Though the ensemble outperforms the other classifiers, it has a signif-  
 327 icant drawback: its training time is significantly larger than that of the individ-  
 328 ual classifiers. Support-vector machines are infamous for their slowness when  
 329 there is a lot of training data, and neural networks can require a lot of training



time whenever the number of neurons gets large (Burges and Schölkopf, 1997;  
Kamathi and Pittner, 1999). Requiring both algorithms to run will therefore  
require a lot of additional training time. Considering the marginal performance  
increase, the cost outweighs the benefit. As a result, when taking everything into  
account, the support-vector machine is the best classifier for labeling conspiracy  
videos on YouTube.

## 5 Discussion

To do.

## 6 Planning

Table 1: Planning

Week	Handelingen	Afgehandeld
28/03-03/04	Hyperparameters optimaliseren	Ja
03/04-10/04	Classifier-ensemble optimaliseren	Ja
11/04-17/04	Google accounts maken, deelvraag 3 maken	Ja
18/04-24/04	Inleiding uitbreiden	Ja
25/04-01/05	Uitvoering experiment, labelen met classifier	
02/05-08/05	Deelvraag 1 schrijven, beginnen deelvraag 2	
09/05-15/05	Deelvraag 2 afschrijven	
16/05-22/05	Beginnen met discussie schrijven	
23/05-29/05	Discussie afschrijven	
30/05-05/06	Abstract schrijven, tekst proof-readen	
06/06-12/06	Laatste aanpassingen en verbeteringen	
13/06-19/06	Inleveren scriptie en verdediging	



## References

- 337 Burges, C. J. and Schölkopf, B. (1997). Improving the accuracy and speed  
338 of support vector machines. *Advances in neural information processing*  
339 *systems*, pages 375–381.
- 340 Cooper, P. (2020). How does the youtube algorithm work? a guide to getting  
341 more views.
- 342 Covington, P., Adams, J., and Sargin, E. (2016). Deep neural networks for  
343 youtube recommendations. In *Proceedings of the 10th ACM conference on*  
344 *recommender systems*, pages 191–198.
- 345 Donzelli, G., Palomba, G., Federigi, I., Aquino, F., Cioni, L., Verani, M., Car-  
346 ducci, A., and Lopalco, P. (2018). Misinformation on vaccination: A quan-  
347 titative analysis of youtube videos. *Human vaccines & immunotherapeutics*,  
348 14(7):1654–1659.
- 349 Feurer, M. and Hutter, F. (2019). Hyperparameter optimization. In *Automated*  
350 *Machine Learning*, pages 3–33. Springer, Cham.
- 351 Kamarathi, S. V. and Pittner, S. (1999). Accelerating neural network training  
352 using weight extrapolations. *Neural networks*, 12(9):1285–1299.
- 353 Karaa, W. B. A. (2013). A new stemmer to improve information retrieval.  
354 *International Journal of Network Security & Its Applications*, 5(4):143.
- 355 Lang, P. (2018). Youtube average view duration - the 50% rule.
- 356 Ledwich, M. and Zaitsev, A. (2019). Algorithmic extremism: Examining  
357 youtube’s rabbit hole of radicalization. *arXiv preprint arXiv:1912.11211*.
- 358 Lemaître, G., Nogueira, F., and Aridas, C. K. (2017). Imbalanced-learn: A  
359 python toolbox to tackle the curse of imbalanced datasets in machine learn-  
360 ing. *The Journal of Machine Learning Research*, 18(1):559–563.
- 361 Lubach, A. (2020). De online fabeltjesfuik — zondag met lubach (s12).
- 362 Maack, M. M. (2019). ‘youtube recommendations are toxic,’ says dev who  
363 worked on the algorithm.
- 364 Miller, D. T. (2021). Characterizing qanon: Analysis of youtube comments  
365 presents new conclusions about a popular conservative conspiracy. *First*  
366 *Monday*.
- 367 Neufeld, D. (2021). The 50 most visited websites in the world.
- 368 Paolillo, J. C. (2018). The flat earth phenomenon on youtube. *First Monday*.
- 369 Pariser, E. (2011). *The filter bubble: What the Internet is hiding from you*.  
370 Penguin UK.



- 371 Park, M., Naaman, M., and Berger, J. (2016). A data-driven study of view  
372 duration on youtube. In *Proceedings of the International AAAI Conference*  
373 *on Web and Social Media*, volume 10.
- 374 Reitermanova, Z. (2010). Data splitting. In *WDS*, volume 10, pages 31–36.
- 375 Rosenbaum, L. (2021). Escaping catch-22—overcoming covid vaccine hesitancy.
- 376 Roth, C., Mazières, A., and Menezes, T. (2020). Tubes and bubbles topological  
377 confinement of youtube recommendations. *PloS one*, 15(4):e0231703.
- 378 Sokolova, M. and Lapalme, G. (2009). A systematic analysis of performance  
379 measures for classification tasks. *Information processing & management*,  
380 45(4):427–437.
- 381 Sun, Y., Kamel, M. S., and Wang, Y. (2006). Boosting for learning multiple  
382 classes with imbalanced class distribution. In *Sixth international conference*  
383 *on data mining (ICDM'06)*, pages 592–602. IEEE.
- 384 YouTube (2021). Youtube community guidelines & policies - how youtube works.
- 385 Zhou, R., Khemmarat, S., and Gao, L. (2010). The impact of youtube rec-  
386 ommendation system on video views. In *Proceedings of the 10th ACM*  
387 *SIGCOMM conference on Internet measurement*, pages 404–410.

# Appendices

## A Hyperparameter tuning

Ensemble	Activation	Layers	Neurons	Accuracy	Precision	Recall	F1
svm, nn, knn	logistic	1	20	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	relu	1	20	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	identity	1	1	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	identity	10	1	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	logistic	1	10	0.939318	0.948923	0.931204	0.93998
ridge, svm, nn, knn	tanh	10	1	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	tanh	1	10	0.939318	0.948923	0.931204	0.93998
svm, nn, knn	tanh	1	20	0.939318	0.948923	0.931204	0.93998
svm, nn, logr, knn	identity	1	1	0.939318	0.948923	0.931204	0.93998
svm, nn, logr, knn	identity	1	10	0.939318	0.948923	0.931204	0.93998

Table 2: Ensemble.

Kernel	C	Accuracy	Precision	Recall	F1
rbf	10.0	0.936309	0.945473	0.928747	0.937035
rbf	100.0	0.935557	0.942289	0.930713	0.936465
rbf	1.0	0.925276	0.930163	0.922850	0.926492
poly	10.0	0.916499	0.946017	0.886978	0.915547
poly	100.0	0.915246	0.944940	0.885504	0.914257
linear	1.0	0.913741	0.917944	0.912531	0.915229
poly	1.0	0.909729	0.935065	0.884521	0.909091
linear	10.0	0.904965	0.907882	0.905651	0.906765
linear	100.0	0.898195	0.905830	0.893366	0.899555
rbf	0.1	0.878887	0.878906	0.884521	0.881705

Table 3: Support-vector machine.



Activation	Layers	Neurons	Accuracy	Precision	Recall	F1
identity	10	10	0.923019	0.935484	0.912039	0.923613
identity	25	10	0.921013	0.933031	0.910565	0.921661
relu	10	10	0.919007	0.917561	0.924324	0.920930
identity	10	20	0.916750	0.906056	0.933661	0.919652
relu	10	20	0.915998	0.912221	0.924324	0.918233
tanh	10	10	0.915747	0.914592	0.920885	0.917728
relu	1	1	0.915747	0.931876	0.900737	0.916042
tanh	25	20	0.915496	0.919052	0.914988	0.917016
tanh	10	20	0.914744	0.920178	0.912039	0.916091
logistic	1	1	0.913741	0.925516	0.903686	0.914470

Table 4: Neural network.

Solver	Alpha	Accuracy	Precision	Recall	F1
auto	0.1	0.918506	0.919118	0.921376	0.920245
sparse_cg	0.1	0.918506	0.919118	0.921376	0.920245
sag	0.1	0.918255	0.919902	0.919902	0.919902
auto	1.0	0.917252	0.923497	0.913514	0.918478
sparse_cg	1.0	0.917252	0.923497	0.913514	0.918478
sag	1.0	0.917252	0.923497	0.913514	0.918478
sag	10.0	0.878385	0.893002	0.865356	0.878962
auto	10.0	0.878134	0.892549	0.865356	0.878743
sparse_cg	10.0	0.878134	0.892549	0.865356	0.878743
auto	100.0	0.812437	0.854545	0.762162	0.805714

Table 5: Ridge regression.

Penalty	C	Solver	Accuracy	Precision	Recall	F1
l2	20	newton-cg	0.918506	0.924107	0.915479	0.919773
l2	20	saga	0.918506	0.924107	0.915479	0.919773
l2	20	sag	0.918506	0.924107	0.915479	0.919773
l2	10	sag	0.916249	0.920831	0.914496	0.917653
l2	10	newton-cg	0.916249	0.920831	0.914496	0.917653
l2	10	saga	0.916249	0.920831	0.914496	0.917653
l2	10	lbfgs	0.915747	0.919506	0.914988	0.917241
l2	20	lbfgs	0.914744	0.921432	0.910565	0.915966
none	1	sag	0.913992	0.923848	0.906143	0.914909
none	10	saga	0.913240	0.922461	0.906143	0.914229

Table 6: Logistic regression.



K	Accuracy	Precision	Recall	F1
1	0.889669	0.888456	0.896314	0.892368
3	0.888415	0.882212	0.901720	0.891859
4	0.879137	0.908899	0.848157	0.877478
5	0.873370	0.858482	0.900246	0.878868
6	0.873119	0.882441	0.866830	0.874566
2	0.872618	0.935043	0.806388	0.865963
7	0.868355	0.848891	0.902703	0.874970
8	0.867603	0.867382	0.874201	0.870778
9	0.861585	0.835672	0.907125	0.869934
10	0.859579	0.854397	0.873710	0.863946

Table 7: K-nearest neighbors.